

Case Study 3: fMRI Prediction

Coping with Large Covariances: Graphical Models, { today Graphical LASSO

Machine Learning for Big Data
CSE547/STAT548, University of Washington

Emily Fox

May 5th, 2015

©Emily Fox 2015

1

Multivariate Normal Models

- So far, we looked at univariate multiple regression

$$y^i = \beta_0 + \beta_1 x_1^i + \dots + \beta_p x_p^i + \epsilon^i \quad \epsilon^i \sim \mathcal{N}(0, \sigma^2) \quad y^i \in \mathbb{R}$$

$$= \beta^T x^i + \epsilon^i$$

$$\Rightarrow y^i \sim \mathcal{N}(\beta^T x^i, \sigma^2)$$

- If one has a multivariate response $y^i \in \mathbb{R}^d$ ← # of semantic features

□ Assuming independence between dimensions

so far

$$y^i \sim \mathcal{N} \left(\begin{bmatrix} \beta^{(1)T} \\ \beta^{(2)T} \\ \vdots \\ \beta^{(d)T} \end{bmatrix} x^i, \begin{bmatrix} \sigma^2 & & 0 \\ & \ddots & \\ 0 & & \sigma^2 \end{bmatrix} \right)$$

$\beta^{(l)}$ are coeff. for the l^{th} semantic feature

leads to ind. problems

©Emily Fox 2015

2

Multivariate Normal Models

- If one has a multivariate response $y^i \in \mathbb{R}^d$

□ Assuming correlation between the output dimensions

$$y^i \sim N(B^T x^i, \Sigma)$$

recall : $\text{cov}(y_s, y_t) = \Sigma_{st}$

- Assume linear (or other mean regression) is removed and focus on the correlation structure

$$y^i \sim N(0, \Sigma) \quad \Sigma \text{ sym. pos. def.}$$

- Matrix valued parameter!

see more on matrix valued params in Case Study 4

©Emily Fox 2015

3

Low-Rank Approximations

- In pictures...

$$\Sigma_0 = \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$$

$$\Sigma = \Lambda \Lambda' + \Sigma_0$$

Diagram illustrating the low-rank approximation of a covariance matrix Σ . The matrix Σ is shown as a large blue square, labeled "Very big matrix" and "d". It is decomposed into a product of a matrix Λ (labeled "d x k") and its transpose Λ' (labeled "k x d"), plus a diagonal matrix Σ_0 (labeled "d x d"). The matrix Λ is shown as a tall blue rectangle, labeled "low rank" and "k". The matrix Σ_0 is shown as a blue square with a diagonal line, labeled "d diag terms" and "maintains pos. def."

- Number of parameters:

$$dk + d = d(k+1) \ll \frac{d(d+1)}{2}$$

sig. reduction in param. for $k \ll d$

©Emily Fox 2015

4

Latent Factor Models

- Original multivariate regression $y^i = B^T x^i + \epsilon^i$, $\epsilon^i \sim N(0, \Sigma)$ here: assume linear term is removed

$$y^i = B^T x^i + \epsilon^i, \quad \epsilon^i \sim N(0, \Sigma)$$

- Latent factor model assumption: $\Sigma = \Lambda \Lambda' + \Sigma_0$
- Low-rank approximation arises from a latent factor model

$$y^i = \Lambda \eta^i + \tilde{\epsilon}^i$$

obs. \uparrow "factor loadings" \uparrow "latent factors" \uparrow ind. across dims

$\eta^i \sim N_k(0, I)$
 $\tilde{\epsilon}^i \sim N_d(0, \Sigma_0)$ \leftarrow diag.

- Proof:


$$\begin{aligned} \text{Cov}(y, \Lambda, \Sigma_0) &= E[(y - E[y])(y - E[y])^T] = E[yy^T] \\ &= E[(\Lambda \eta + \tilde{\epsilon})(\Lambda \eta + \tilde{\epsilon})^T] = \Lambda E[\eta \eta^T] \Lambda + 2E[\eta \tilde{\epsilon}] \Lambda^T + E[\tilde{\epsilon} \tilde{\epsilon}^T] \\ &= \Lambda I \Lambda^T + \Sigma_0 \end{aligned}$$

5

©Emily Fox 2015

Lower-dim Embeddings

Sharing information in *low-dim subspace*



©Emily Fox 2015

6

Sparsity Assumptions

- What if we assume Σ is sparse?

$$(i \neq j) \Sigma_{ij} = 0 \xRightarrow{\text{Gaussian}} y_i \perp\!\!\!\perp y_j \quad \leftarrow \text{ind.}$$

$$\text{cov}(y_i, y_j) = 0$$

Could assume Σ sparse to reduce # params,
but each 0 encodes an indep. statement
 \rightarrow often too strong of an assumption

- More often, we can reasonably make statements about conditional independence

$$\text{"cat"} \perp\!\!\!\perp \text{"dog"} \mid \text{"animal", "furry", "pet"} \quad \leftarrow \text{cond. on}$$

©Emily Fox 2015

7

Information Form

- Motivations for considering "information form" of multivariate normal

- Easier to read off conditional densities
- Has log-linear form in terms of "information parameters"

$$\frac{1}{\sqrt{2\pi} |\Sigma|} e^{-\frac{1}{2}(y-\mu)^T \Sigma^{-1} (y-\mu)} \quad \leftarrow y \sim N(\mu, \Sigma)$$

$$\begin{aligned} \uparrow \Omega &= \Sigma^{-1} \\ \downarrow \eta &= \Sigma^{-1} \mu \end{aligned}$$

$$\propto e^{\eta^T y - \frac{1}{2} y^T \Omega y}$$

$$\begin{aligned} & y^T \Sigma^{-1} y \\ & - 2 y^T \Sigma^{-1} \mu \\ & + \mu^T \Sigma^{-1} \mu \end{aligned} \quad \leftarrow \text{const. wrt } y$$

$$\leftarrow y \sim N^{-1}(\eta, \Omega)$$


©Emily Fox 2015

8

Conditional Densities

- Assume a model with

$$y \sim N^{-1}(\eta, \Omega)$$



and divide the dimensions into two sets A, \bar{A}

- Then,

$$\begin{bmatrix} y_A \\ y_{\bar{A}} \end{bmatrix} \sim N \left(\begin{bmatrix} \eta_A \\ \eta_{\bar{A}} \end{bmatrix}, \begin{bmatrix} \Omega_{AA} & \Omega_{A\bar{A}} \\ \Omega_{\bar{A}A} & \Omega_{\bar{A}\bar{A}} \end{bmatrix} \right)$$

Submatrix of Ω with rows indexed by indices in A and columns by \bar{A}

$$p(y_A | y_{\bar{A}}) = N^{-1}(\eta_A - \Omega_{A\bar{A}} y_{\bar{A}}, \underline{\underline{\Omega_{AA}}})$$

©Emily Fox 2015


9

Conditional Densities

- Let $A = \{s, t\}$

"cat" "dog"

\bar{A} = everything else



y_s, y_t y_{-st}

$$p(y_A | y_{\bar{A}}) = N^{-1}(\eta_A - \Omega_{A\bar{A}} y_{\bar{A}}, \Omega_{AA})$$

what if $\Omega_{st} = 0$? $\Rightarrow \begin{bmatrix} \Omega_{ss} & 0 \\ 0 & \Omega_{tt} \end{bmatrix}$

$$\begin{bmatrix} \Omega_{ss} & \Omega_{st} \\ \Omega_{ts} & \Omega_{tt} \end{bmatrix}$$

$$\text{cov}(y_s, y_t | y_{-st}) = \Omega_{AA}^{-1} = \begin{bmatrix} \Omega_{ss}^{-1} & 0 \\ 0 & \Omega_{tt}^{-1} \end{bmatrix}$$

- Therefore,


$$\Leftrightarrow \boxed{y_s \perp\!\!\!\perp y_t | y_{-st} \Leftrightarrow \Omega_{st} = 0}$$

©Emily Fox 2015

10

Connection with Graphical Models

- Undirected graphical model or Markov random field (MRF)



$$p(y | \eta, \Omega) \propto \prod_t \psi_t(y_t) \prod_{(s,t) \in E} \psi_{st}(y_s, y_t) \quad \left. \begin{array}{l} \psi_t(y_t) \propto e^{\eta_t y_t} \\ \psi_{st}(y_s, y_t) \propto e^{-\frac{1}{2} y_s \Omega_{st} y_t} \end{array} \right\}$$

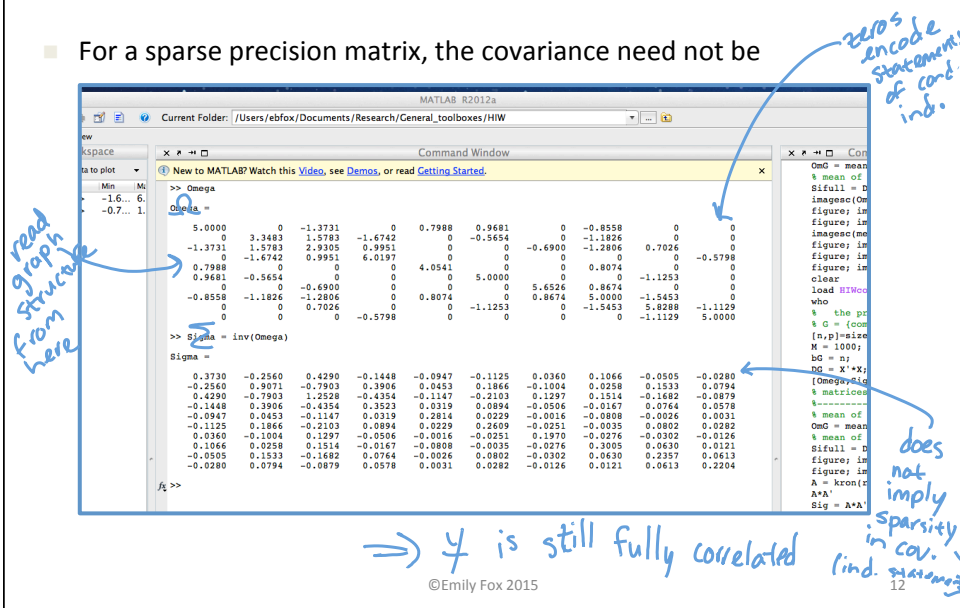
node potentials edge potentials

©Emily Fox 2015

11

Sparse Precision vs. Covariance

- For a sparse precision matrix, the covariance need not be



ML Estimation for Given Graph

- Assume a known graph $G = \{V, E\}$ *nodes edges*
- Rewrite log likelihood:

$$\log p(y|\theta) = \frac{N}{2} \log |\Omega| - \frac{1}{2} \sum_i (y_i - \mu)^T \Omega (y_i - \mu)$$

$$= \frac{N}{2} \log |\Omega| - \frac{1}{2} \sum_i \text{tr} \left[(y_i - \mu)(y_i - \mu)^T \Omega \right]$$


$$\triangleq \frac{N}{2} \log |\Omega| - \frac{1}{2} \text{tr} (S_\mu \Omega)$$

$$S_\mu = \sum_i (y_i - \mu)(y_i - \mu)^T$$

$$L(\Omega) = \log |\Omega| - \text{tr}(S \Omega)$$

In our case, $\mu = 0$

$$\frac{1}{\sqrt{2\pi}|\Sigma|} e^{-\frac{1}{2}(y-\mu)^T \Sigma^{-1}(y-\mu)}$$



trace trick:
 $X^T A X = \text{tr}(X^T A X)$
 $= \text{tr}(X X^T A)$

©Emily Fox 2015

13

ML Estimation for Given Graph


$$L(\Omega) = \log |\Omega| - \text{tr}(S \Omega)$$

- Take gradient:

$$\nabla L(\Omega) = \Omega^{-1} - S$$

$$\text{s.t. } \Omega_{st} = 0 \text{ if } (s, t) \notin E$$

Ω pos. def., sym. matrix



linear constraint

hard

- Many approaches to solving:

- Barrier method – add penalty discouraging Ω from leaving the positive definite cone (Dahl et al. 2008)
- Coordinate descent method (cf., Hastie et al. 2009)
- ...

©Emily Fox 2015

14

ML Estimation for Given Graph

- Can show that the optimal solution satisfies

$$\hat{\Sigma}_{st}^{ML,G} = S_{st} \quad \begin{array}{l} \text{if } (s,t) \in E \\ \text{if } s=t \end{array} \quad \begin{array}{l} \text{match sample} \\ \text{cov.} \end{array}$$

$$\hat{\Omega}_{st} = 0 \quad \text{if } (s,t) \notin E$$

- Example:

adj. matrix
1 = edge

$$G = \begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 0 \end{pmatrix} \quad S = \begin{pmatrix} 10 & 1 & 5 & 4 \\ 1 & 10 & 2 & 6 \\ 5 & 2 & 10 & 3 \\ 4 & 6 & 3 & 10 \end{pmatrix}$$

$$\hat{\Sigma}^{-1} = \Omega = \begin{pmatrix} \cdot & \cdot & 0 & \cdot \\ \cdot & \cdot & \cdot & 0 \\ 0 & \cdot & \cdot & \cdot \\ \cdot & 0 & \cdot & \cdot \end{pmatrix} \quad \hat{\Sigma}^{ML,G} = \Sigma = \begin{pmatrix} 10 & 1 & 1.31 & 4 \\ 1 & 10 & 2 & 0.87 \\ 1.31 & 2 & 10 & 3 \\ 4 & 0.87 & 3 & 10 \end{pmatrix}$$

©Emily Fox 2015

15

Estimating Graph Structure

- To learn the structure of the Gaussian graphical model, we want to trade off fit and sparsity

□ Measure of fit: log-likelihood

$$\log |\Omega| - \text{tr}(S\Omega) + \text{const.}$$


□ Encouraging sparsity: $\Omega_{st} = 0 \Rightarrow$ no edge "sparsity"

$$\|\Omega\|_1 = \sum_{s,t} |\Omega_{s,t}| \quad \leftarrow \text{want to min}$$

- Overall objective = "graphical LASSO" or "Glasso"

$$F(\Omega) = -\log |\Omega| + \text{tr}(S\Omega) + \lambda \|\Omega\|_1$$

Just as in LASSO, but w/ a matrix parameter and s.t. $\Omega \succ 0$



©Emily Fox 2015

16

Solving the Graphical LASSO

- Objective is convex, but non-smooth as in LASSO ... *subgrad.*
- Also, positive definite constraint!
- There are many approaches to optimizing the objective
 - Most common = coordinate descent akin to shooting algorithm (Friedman et al. 2008) *see HW 3*
- Some issues...
 - Ballpark: several minutes for a 1000-variable problem
 - Algorithms scale as $O(d^3)$
- Other approach = ADMM *also HW 3*

©Emily Fox 2015

17

Faster Computations

From Daniela Witten's talk at JSM 2012:

1. The j th variable is unconnected from all others in the graphical lasso solution if and only if $|S_{ij}| \leq \lambda$ for all $i = 1, \dots, j-1, j+1, \dots, p$. *Sample cov is small relative to chosen penalty*
 2. Let \mathbf{A} denote the $p \times p$ matrix whose elements take the form $A_{ij} = 1, A_{ij} = 1_{|S_{ij}| > \lambda}$. Then the connected components of \mathbf{A} are the same as the connected components of the graphical lasso solution. *ind. on the thresholded values*
- We can obtain the exact right answer by solving the graphical lasso on each connected component separately!


Citations: Witten et al. JCGS 2011, Mazumder and Hastie JMLR 2012

©Emily Fox 2015

18

Covariance Screening for Glasso

From Daniela Witten's talk at JSM 2012:



- ▶ The solution to the graphical lasso problem with $\lambda = 0.7$ has five connected components (why 5?!)
- ▶ Perform graphical lasso on each component separately!
- ▶ **Reduction in computational time:** From $O(50^3)$ to $O(24^3)$.