Case Study 5: Mixed Membership Modeling

Variational Inference for LDA

Machine Learning for Big Data
CSE547/STAT548, University of Washington
Emily Fox
June 2nd, 2015

Variational Methods Goal

- Recall task: Characterize the posterior $p(\theta, z \mid x)$ (obs.; params; latent vars)

- Turn posterior inference into an optimization task
- Introduce a "tractable" family of distributions over parameters and latent variables
  - Family is indexed by a set of "free parameters"
  - Find member of the family closest to: $p(\theta, z \mid x)$ (want q closest to p; all 2D Gaussians)

Call family Q and want $q \in Q$ that is closest to $p(\theta, z \mid x)$

Variational Methods Cartoon

- Cartoon of goal:

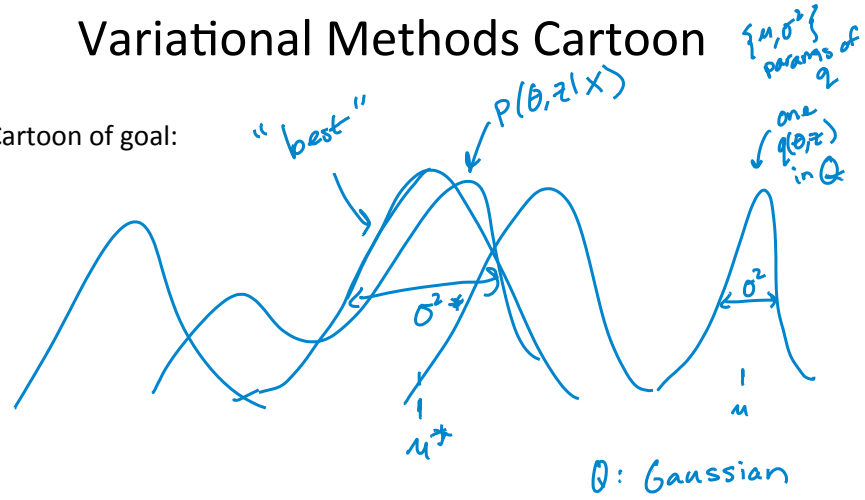

- Questions:
  1. How do we measure "closeness"?
  2. If the posterior is intractable, how can we approximate something we do not have to begin with?

Interpretations of Minimizing Reverse KL

- Evidence lower bound (ELBO)

KL div. = measure of "distance" between p + q

$$\log p(x) = D(q(z,\theta)||p(z,\theta|x)) + \mathcal{L}(q) \geq \mathcal{L}(q)$$

const. — add to a const. — ELBO

- Therefore,
  - ELBO provides a lower bound on marginal likelihood
  - Maximizing ELBO is equivalent to minimizing KL

$$\max \mathcal{L}(q) = \min D(q\|p) = \max \text{ lower bound on } \log p(x)$$

what we can control — depends on what we don't know

Mean Field

ELBO $\mathcal{L}(q) = E_q[\log p(z, \theta, x)] - E_q[\log q(z, \theta)]$

- How do we choose a Q such that the following is tractable?

$\hat{q} = \arg\max_{q \in Q} \mathcal{L}(q)$ ← new objective

- Simplest case = mean field approximation ← to make objective tractable = how we choose structure of Q
  - Assume each parameter and latent variable is conditionally independent given the set of free parameters

$\theta, z = \{z^1, \ldots, z^N\}$

$q(\theta, z) = q(\theta|\gamma) \prod_{i=1}^{N} q(z^i|\phi^i)$ ("free param")

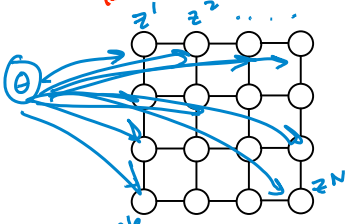

Original graph

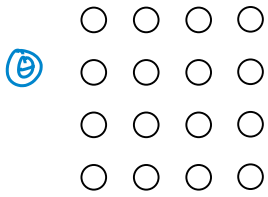

Naïve mean field

example

our focus, but there also is structured approx.

Mean Field for LDA

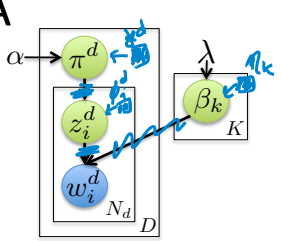

- In LDA, our parameters are $\theta = \{\pi^d\}, \{\beta_k\}$
  $z = \{z_i^d\}$
- The variational distribution factorizes as

$q(\pi, \beta, z) = \prod_{k=1}^{K} q(\beta_k|\eta_k) \prod_{d=1}^{D} \left[q(\pi^d|\gamma^d) \prod_{i=1}^{N_d} q(z_i^d|\phi_i^d) \right]$

$\{\pi^1, \ldots, \pi^D\}$

$q(\beta_k|\eta_k)$: $\text{Dir}(\eta_{k1}, \ldots, \eta_{kV})$

$q(\pi^d|\gamma^d)$: $\text{Dir}(\gamma_1^d, \ldots, \gamma_K^d)$

$q(z_i^d|\phi_i^d)$: $\text{Mult}(\phi_i^d)$, $\sum_{k=1}^{K} \phi_{ik}^d = 1$, $i \in 1, \ldots, K$

optimize these free params

need to enforce this constraint

- The joint distribution factorizes as

$$p(\pi, \beta, z, w) = \prod_{k=1}^{K} p(\beta_k \mid \lambda) \prod_{d=1}^{D} p(\pi^d \mid \alpha) \prod_{i=1}^{N_d} p(z_i^d \mid \pi^d) p(w_i^d \mid z_i^d, \beta)$$

Mean Field for LDA

$$q(\pi, \beta, z) = \prod_{k=1}^{K} q(\beta_k \mid \eta_k) \prod_{d=1}^{D} q(\pi^d \mid \gamma^d) \prod_{i=1}^{N_d} q(z_i^d \mid \phi_i^d)$$

mean field approx.

$$p(\pi, \beta, z, w) = \prod_{k=1}^{K} p(\beta_k \mid \lambda) \prod_{d=1}^{D} p(\pi^d \mid \alpha) \prod_{i=1}^{N_d} p(z_i^d \mid \pi^d) p(w_i^d \mid z_i^d, \beta)$$

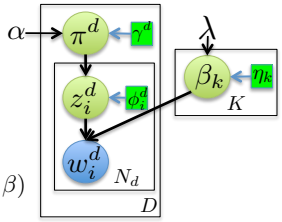

- Examine the ELBO

$$\mathcal{L}(q) = \sum_{k=1}^{K} E_q[\log p(\beta_k \mid \lambda)] + \sum_{d=1}^{D} E_q[\log p(\pi^d \mid \alpha)] + \sum_{d=1}^{d} \sum_{i=1}^{N_d} E_q[\log p(z_i^d \mid \pi^d)] + E_q[\log p(w_i^d \mid z_i^d, \beta)]$$

from joint

$$- \sum_{k=1}^{K} E_q[\log q(\beta_k \mid \eta_k)] - \sum_{d=1}^{D} E_q[\log q(\pi^d \mid \gamma^d)] - \sum_{d=1}^{d} \sum_{i=1}^{N_d} E_q[\log q(z_i^d \mid \phi_i^d)]$$

all terms from q

Optimize via Coordinate Ascent

- Algorithm:

per doc params

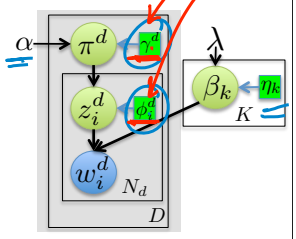

for $d = 1, \ldots, D$

$$\frac{\partial \mathcal{L}}{\partial \gamma^d} = 0 \rightarrow \gamma^{d(t+1)} = \alpha + \sum_{i=1}^{N_d} \phi_i^{d(t)}$$

for $i = 1, \ldots, N_d$

pmf over topics $1, \ldots, K$

$$\frac{\partial \mathcal{L}}{\partial \phi_i^d} = 0 \rightarrow \phi_i^{d(t+1)} \propto \exp\left\{\Psi(\gamma_{1:K}^{d(t+1)}) + \Psi(\eta_{1:K,w_i^d}^{(t)}) - \Psi(\textstyle\sum \eta_{1:K,w}^{(t)})\right\}$$

use Lagrange multiplier to enforce $\sum \phi_i^d = 1$

DATA PARALLEL across ϕ_i^d, γ^d

Optimize via Coordinate Ascent

- Algorithm:

for $k = 1, \ldots, K$

$$\frac{\partial \mathcal{L}}{\partial \eta_k} = 0 \rightarrow \eta_k^{(t+1)} = \lambda + \sum_{d=1}^{D} \sum_{i=1}^{N_d} w_i^d \phi_i^{d(t+1)}$$

aggregate

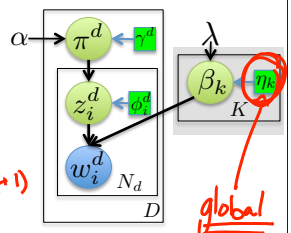

global free param (touches every doc)

Map Reduce

Generalizing

- Many Bayesian model have this form:

$$p(\theta, z^{1:N}, x^{1:N}) = p(\theta) \prod_{i=1}^{N} p(z^i \mid \theta) p(x^i \mid z^i, \theta)$$

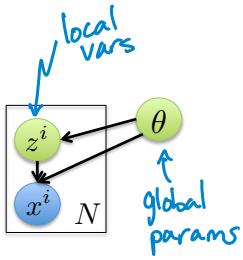

- Goal is to compute $p(\theta, z^{1:N} \mid x^{1:N})$

- Assume each complete conditional is in the exponential family

$$p(z^i \mid \theta, x^i) = h(z^i) \exp\{\eta_\ell(\theta, x^i)^T z^i - a(\eta_\ell(\theta, x^i))\}$$

"local"

$$p(\theta \mid z, x) = h(\theta) \exp\{\eta_g(z, x)^T \theta - a(\eta_g(z, x))\}$$

global

Generalizing

- Mean field variational approximation

$$q(z, \theta) = q(\theta \mid \gamma) \prod_{i=1}^{N} q(z^i \mid \phi^i)$$

- Match each component to have same family as model conditional

e.g. for global params

$$p(\theta \mid z, x) = h(\theta) \exp\{\eta_g(z, x)^T \theta - a(\eta_g(z, x))\}$$

$$q(\theta \mid \gamma) = h(\theta) \exp\{\gamma^T \theta - a(\gamma)\}$$

- Same for local variational terms, too

Generalizing

- Under these exponential family assumptions, the gradient is:

$$\nabla_\gamma \mathcal{L} = a''(\gamma)(E_\phi[\eta_g(z, x)] - \gamma)$$

- This leads to a simple coordinate update (Ghahramani and Beal, 2001)

$$\nabla_\gamma \mathcal{L} = 0 \rightarrow \gamma^* = E_\phi[\eta_g(z, x)]$$

General Coord. Ascent Algorithm

Initialize γ randomly.
Repeat until the ELBO converges

1. For each data point, update the local variational parameters:

$$\phi_i^{(t)} = \mathrm{E}_{\gamma^{(t-1)}}[\eta_\ell(\theta, x_i)] \text{ for } i \in \{1, \ldots N\}.$$

2. Update the global variational parameters:

$$\gamma^{(t)} = \mathrm{E}_{\phi^{(t)}}[\eta_g(Z_{1:N}, x_{1:N})]$$

Case Study 5: Mixed Membership Modeling

Stochastic Variational Inference

Machine Learning for Big Data
CSE547/STAT548, University of Washington
Emily Fox
June 2nd, 2015

Limitations of Batch Variational Methods

interested in huge datasets (e.g. millions of docs)

Limitations of Batch Variational Methods

- *Example = LDA*
  - Start from randomly initialized η_k (topics)
  - Analyze whole corpus before updating η_k again
- Streaming data: can't compute one iteration!

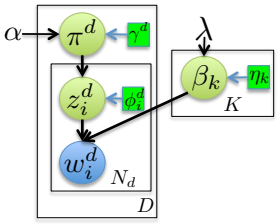

- *More generally…*
  - Do some local computation for each data point.
  - Aggregate these computations to re-estimate global structure.
  - Repeat.
- Inefficient, and cannot handle massive data sets.

Stochastic Variational Inference

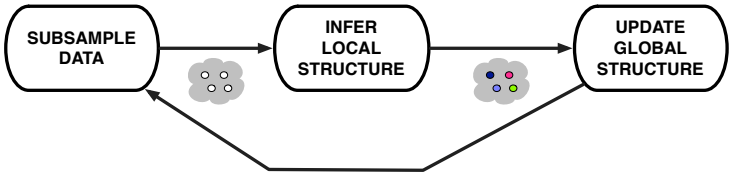

- Stochastic variational inference harnesses:
  - Idea #1: **Stochastic optimization** (Robbins and Monro, 1951)
  - Idea #2: **Natural gradients** (Amari, 1998)

same ideas as behind SGD

Alternative Optimization Schemes

- Didn't have to do coord. ascent. Could have used gradient ascent.

$\hookrightarrow \nabla_\gamma \mathcal{L} = 0$

$\gamma^{(t+1)} = \gamma^{(t)} + \epsilon_t \nabla_\gamma \mathcal{L}$

- Here,

$$\mathcal{L}(q) = E_q[\log p(\theta)] - E_q[\log q(\theta)] + \sum_{i=1}^{N} E_q[\log p(z^i, x^i \mid \theta)] - E_q[\log q(z^i)]$$

touches all of the data (docs)!

- Use **stochastic gradient** step instead
  - Consider one data point x^t sampled uniformly at random and define:

$$\mathcal{L}_t(q) = E_q[\log p(\theta)] - E_q[\log q(\theta)] + N(E_q[\log p(z^t, x^t \mid \theta)] - E_q[\log q(z^t)])$$

"t-ELBO"

$E[\nabla \mathcal{L}_t] = \nabla \mathcal{L}$ ← using all data

over subsampled data — using subsample

Alternative Optimization Schemes

- Recall the gradient of the ELBO for the global parameter:

$$\nabla_\gamma \mathcal{L} = a''(\gamma)(E_\phi[\eta_g(z,x)] - \gamma)$$

↑ global free param

- Even using just one data point, issue for scalability:

must compute $a''(\gamma)$...

computationally intensive

Natural Gradient of the ELBO

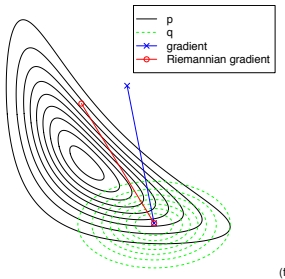

(from Honkela et al., 2010)

doubly beneficial here

① takes geometry into account

② removes $a''(\gamma)$

- The **natural gradient** accounts for the geometry of parameter space
- Natural gradient of the ELBO:

$$\hat{\nabla}_\gamma \mathcal{L} = E_\phi[\eta_g(z,x)] - \gamma$$

↑ natural gradient ↖ "global"

Noisy Natural Gradients

- Let $\eta_t(z^t, x^t)$ be the conditional distribution of the global parameters for the model where the observations are *N* replicates of x^t
- With this, the noisy natural gradient of the ELBO is

$$\hat{\nabla}_\gamma \mathcal{L}_t = \mathrm{E}_{\phi_t}[\eta_t(z^t, x^t)] - \gamma$$

↑ natural grad of t-ELBO

- Notes:
  - It only requires the local variational parameters of one data point.
  - In contrast, the full natural gradient requires all local parameters.
  - Thanks to conjugacy it has a simple form.

SVI Algorithm Overview

Initialize global parameters $\hat{\gamma}$ randomly.
Set the step-size schedule ϵ_t appropriately.
Repeat forever

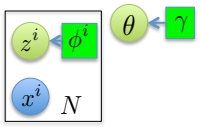

1. Sample a data point uniformly,

$$x_t \sim \mathrm{Uniform}(x_1, \ldots, x_N).$$

2. Compute its local variational parameter,

$$\phi_t = \mathrm{E}_{\gamma^{(t-1)}}[\eta_l(\theta, x_t)]$$

just as in coord. asc., but only for local var

3. Pretend its the only data point in the data set,

$$\hat{\gamma} = \mathrm{E}_{\phi_t}[\eta_t(Z_t, x_t)]$$

4. Update the current global variational parameter,

$$\gamma^{(t)} = (1 - \epsilon_t)\gamma^{(t-1)} + \epsilon_t \hat{\gamma}.$$

$\hat{\nabla}_\gamma \mathcal{L}_t = \mathrm{E}_\phi[\cdot] - \gamma$

$\gamma^{(t)} = \gamma^{(t-1)} + \epsilon_t \hat{\nabla}_\gamma \mathcal{L}_t$
$= \gamma^{(t-1)} + \epsilon_t \mathrm{E}_{\phi_t}[\cdot] - \epsilon_t \gamma^{(t-1)}$
$\hat{\gamma}$

SVI for LDA

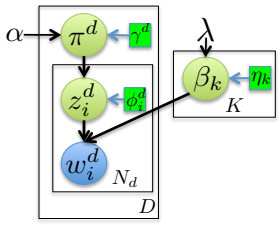

- In LDA, the full ELBO is given by

all docs

$$\mathcal{L} = E_q[\log p(\beta)] - E_q[\log q(\beta)] + \sum_{d=1}^{D} E_q[\log p(\pi^d)] - E_q[\log q(\pi^d)]$$
$$+ \sum_{d=1}^{D} E_q[\log p(z^d, x^d \mid \pi^d, \beta)] - E_q[\log q(z^d)]$$

ELBO

- Assuming *D* documents, consider one sampled at random

act as if doc t is repeated D times

$$\mathcal{L}_t = E_q[\log p(\beta)] - E_q[\log q(\beta)] + D\left(E_q[\log p(\pi^t)] - E[\log q(\pi^t)]\right)$$
$$+ D\left(E_q[\log p(z^t, x^t \mid \pi^t, \beta)] - E_q[\log q(z^t)]\right)$$

t-ELBO

just doc t

SVI for LDA

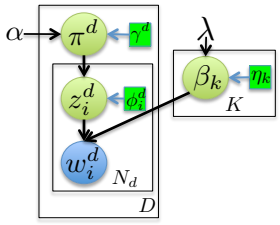

- Initialize $\eta^{(0)}$ randomly.
- Repeat (indefinitely):
  - Sample a document ***d*** uniformly from the data set.
  - For all *k*, initialize $\gamma_k^d = 1$
  - Repeat ~~until converged~~
    - For i=1,…,N_d

local free params for doc d

$$\phi_{ik}^d \propto \exp\{E[\log \pi_k^d] + E[\log \beta_{k,w_i^d}]\}$$

just as in coord. asc. for this doc

  - Set

$$\gamma^d = \alpha + \sum_{i=1}^{N_d} \phi_i^d$$

  - Take a stochastic gradient step $\eta^{(t)} = \eta^{(t-1)} + \epsilon_t \nabla_\eta \mathcal{L}_d$

global vars

$$\lambda + D \sum_{i=1}^{N_d} \phi_i^d w_i^d - \eta^{(t-1)}$$

looks exactly like coord. asc. update for doc d replicated D times

$$\eta^{(t)} = (1-\epsilon_t)\eta^{(t-1)} + \epsilon_t (\quad)$$

SVI for LDA in Practice

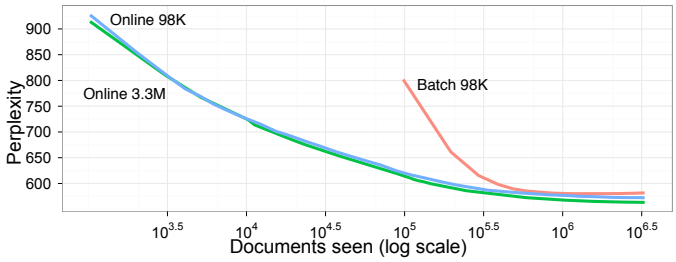

Documents analyzed	2048	4096	8192	12288	16384	32768	49152	65536
Top eight words	systems road made service announced national west language	systems health communication service billion language care road	service systems health companies market communication company billion	service systems companies business company billion health industry	service companies systems business company industry market billion	business service companies industry company management systems services	business service companies industry services company management public	business industry service companies services company management public

(Hoffman et al. 2010)

What you need to know…

- Variational methods
  - Overall goal
  - Interpretation in terms of minimizing (reverse) KL
  - Mean field approximation
- Mean field variational inference for LDA
- Stochastic variational inference
  - General idea of using natural gradients + stochastic optimization
  - Resulting generic algorithm
  - SVI for LDA

Reading

- **Inference in LDA:**
  - Basic LDA and batch variational inference in LDA: Blei, David M., Andrew Y. Ng, and Michael I. Jordan. "Latent dirichlet allocation." the Journal of machine Learning research 3 (2003): 993-1022.
  - Stochastic variational inference: Hoffman, Matt, et al. "Stochastic Variational Inference." arXiv:1206.7051 (2012).

Acknowledgements

- Thanks to Dave Blei for some material in this lecture relating to SVI

Course Wrapup

Overview of CSE 547 / STAT 548 Topics Covered

Machine Learning for Big Data
CSE547/STAT548, University of Washington
Emily Fox
June 2nd, 2015

What you need to know

- Case Study 1: Estimating Click Probabilities
  - Logistic regression
  - Regularization
  - Gradient descent, stochastic gradient decent
  - Hashing and sketching, random projections

What you need to know

- Case Study 2: Document Retrieval and Clustering
  - Approach 1: **k-NN for nearest neighbor search**
  - *Algorithm:* Fast k-NN using KD-trees (exact)
  - *Algorithm:* Approximate k-NN using KD-trees and locality sensitive hashing
  - Approach 2: **k-means for clustering**
  - Data parallel problems
  - *Algorithm*: MapReduce framework and parallel k-means using MapReduce
  - Approach 3: **Gaussian mixture models (GMM) for clustering**
  - *Algorithm*: EM

What you need to know

- Case Study 3: fMRI Prediction
  - Regularized linear models: Ridge regression and LASSO
  - Sparsistency
  - LASSO solvers:
    - LARS
    - Shotgun (stochastic coordinate descent)
    - Hogwild (stochastic gradient descent)
    - Averaging methods
    - ADMM
  - LASSO variants:
    - Fused LASSO
    - Graphical LASSO
  - Coping with large covariances using latent factor models

What you need to know

- Case Study 4: Collaborative Filtering
  - Approach: Matrix factorization
  - *Algorithm*: Alternating least squares (ALS)
  - *Algorithm*: Stochastic gradient descent (SGD)
  - Cold-start problem and feature-based collaborative filtering
  - Model variants:
    - Non-negative matrix factorization
    - Probabilistic matrix factorization
      - *Algorithm*: Gibbs sampling
    - Probabilistic latent space models and network data
  - Graph parallel problems
  - GraphLab framework and application to distributed ALS and Gibbs for matrix factorization

What you need to know

- Case Study 5: Document Mixed Membership Modeling
  - Approach 1: Bayesian document clustering model
  - Conditional independencies in directed graphical models
  - *Algorithm*: Gibbs sampling and collapsed Gibbs sampling
  - Approach 2: Latent Dirichlet allocation (LDA)
  - *Algorithm*: Collapsed Gibbs sampling
  - *Algorithm*: Variational methods and stochastic variational inference

THANK YOU!!!

- You have been a great, interactive class! …especially for a 9:30am lecture =)
- We're looking forward to the poster session

- Thanks to Marco and Alden, too!

